



CAPTION STUDIO

Commercial Product Audit

Android-first • local captions • unrestricted creative control

DECISION

Proceed with a focused Android beta. Win on local privacy, responsive direct manipulation, one unified font system, and no artificial limits. Do not position v1 as a full CapCut replacement.

Executive verdict

Caption Studio has a credible wedge in a crowded creator market: make the most common caption workflow dramatically clearer and less extractive. The opportunity is not “another video editor.” It is the fastest way to import a clip, generate private on-device captions, correct them, style them precisely, and export without credits, a watermark, or an API key.

- Ship the editor before broad video effects. The subtitle experience is the product; video editing is supporting infrastructure.
- Keep the free baseline genuinely useful: local transcription, unlimited projects/presets/fonts, SRT, and watermark-free video export.
- Monetize optional convenience and expansion—curated asset packs, advanced batch workflows, cloud backup or team features—not access to users’ own fonts, minutes, or exports.
- Treat touch responsiveness as a release gate. Direct manipulation must feel immediate, forgiving, reversible, and visibly tied to the object being changed.

Current build readiness

Area	Status	Commercial meaning
Import + preview	Working beta	A real video can enter the editor; source aspect is retained.
Local transcription	Integrated	Native audio extraction + whisper.rn; model downloads once for offline use.
Caption editing	Working beta	Timeline blocks, text edit, timing handles, delete, style scope.
Direct manipulation	Overhauled	One-finger move, two-finger scale/rotate, large edge handles; requires Seeker acceptance test.
Video transforms	Overhauled	Fit/fill, scale, pan, 90° and free rotation; requires Seeker acceptance test.
Burned-in MP4 export	Not complete	Primary launch blocker; current beta is an editor/transcription milestone, not store-ready.

Market and customer

Primary job to be done

CORE PROMISE

Import a video, generate accurate captions locally, and control how every caption looks and behaves—without hunting, credits, watermarks, or cloud upload.

Priority segments

Segment	Pain today	What wins
Short-form creators	Fast publishing but repetitive subscription costs and fragmented style menus.	Two-tap generation, TikTok-safe canvases, reusable styles, fast corrections.
Privacy-sensitive creators	Interviews, client drafts, school or workplace media cannot casually leave-device.	A credible local-only workflow and clear privacy status.
Budget creators	Credits, export quotas, watermarks, and paywalled fonts turn a basic workflow into recurring spend.	No artificial usage limits and transparent optional purchases.
Power caption stylists	Global edits destroy intentional exceptions; word emphasis is tedious.	Project → caption → word inheritance and explicit apply scope.
Android-first users	Some AI creator tools favor iOS or cloud-first experiences.	Good Android performance, sideloadable beta, broad modern device support.

Competitive signals

The category validates willingness to pay, but it is fragmented between broad editors, template apps, and AI-first services. Published plan pages change frequently, so prices below are reference points checked in August 2026, not permanent guarantees.

Product	Public commercial signal	Gap Caption Studio can own
Mojo	Template-led editor; Mojo Pro page advertises recurring plans, while Pro AI adds credits and additional AI tools.	One coherent caption workspace, unified fonts, explicit one/all scope, local processing.
CapCut	Broad editor with Standard/Pro tiers and increasing feature prompts around paid plans.	A calmer caption-first workflow and predictable ownership.
Captions	Free, Pro, Max and Scale plans; AI creation and higher-level workflows drive upgrades.	Android-first local caption editing without account or cloud dependency.
Zeemo / peers	Caption and translation services commonly compete on AI features and recurring access.	Privacy, offline reliability, tactile editing, and no minute economy.

Evidence links: [Mojo Pro](#) • [Mojo Pro AI](#) • [CapCut pricing](#) • [Captions plans](#)

Mojo recording audit

The supplied 6:36 portrait screen recording was reviewed as a workflow reference. It shows Mojo Free on the Seeker navigating templates, media import, caption generation, progress, caption styling, timeline/transcript controls, fonts and gated features. It has no audio, so the audit is based on visible behavior only.

What the flow does

1. Starts from a broad template-first home, then enters a project and imports media.
2. Launches automatic captions through a guided generation flow with visible progress.
3. Places generated text over a live preview and offers a timeline/transcript editing surface.
4. Provides caption styles, colors, fonts and motion options, while distributing font-related choices across multiple paths and placing some capabilities behind Pro.

KEEP / IMPROVE / IGNORE / LATER

Class	Observed behavior	Caption Studio response
KEEP	Import → generate; progress state; live preview; playhead timeline; transcript list; word chips; visual presets; canvas transforms.	Retain the speed and immediacy. Make every job cancellable and recoverable.
IMPROVE	Style scope is ambiguous; fonts appear in multiple contexts; intro/template framing is heavy; timeline lanes are not optimized for precise correction.	Ask This subtitle / All subtitles; route every font entry to one browser; begin from projects; enlarge timing and transform targets.
IGNORE	Template-first discovery, generic creator dashboards, forced intro framing, and a social-template labyrinth.	Do not copy them into the core caption flow.
LATER	Translation, pause removal, intros, overlays, general effects and broad editing tools.	Add only after transcription, correction, styling and export are excellent.

Detailed UX findings

Area	Finding	Decision
Generation	Clear enough to understand, but can feel like a feature inside a larger template product.	One primary Generate captions action after import.
Timeline	Visual blocks help navigation; precision operations need larger touch targets and clearer state.	Pause + seek on selection, distinct selected styling, start/end grips, delete nearby, undo.
Text edit	Transcript-oriented correction is useful.	Keep block text editing; add word selection later without changing project schema.
Fonts	Font access is fragmented across styles/templates/packs.	One searchable master library: favorites, recent, built-in, imported TTF/OTF.
Color/style	Visual presets accelerate first use.	Keep presets, expose resolved values, and preserve per-caption exceptions.
Animations	Useful for creator content but easy to over-prioritize.	Data-driven presets after stable timing/export.
Pro gates	Visible locks make the exploration feel commercially constrained.	Never gate user fonts, caption minutes, presets, resolutions, watermarks or export count.

Exact v1 experience

Screen map

Screen	Purpose	Primary actions
Projects	Resume or start work without template noise.	New project, import video, recent projects, storage/model status.
Import	Confirm source and canvas before transcription.	Original, 9:16, 16:9, 1:1, 4:5; Fit/Fill; generate captions.
Transcription job	Make a long local job understandable.	Fast/Balanced/Accurate, download state, stage progress, cancel/background, retry.
Editor	Single home for preview + timeline + tools.	Captions, Style, Fonts, Animation, Video, Audio, Export.
Caption inspector	Correct one selected block.	Edit text, timing, split, merge, duplicate, delete, reposition.
Apply scope sheet	Prevent silent style mistakes.	This subtitle, All subtitles; later Same style.
Master fonts	Every font in one predictable place.	Search, favorites, recent, built-in, imported, live sample, import TTF/OTF.
Export	Create a portable result.	MP4, SRT, ASS; 720p, 1080p, Original; progress, cancel, save/share.

Direct manipulation contract

- Tap selects and shows a high-contrast bounding box. Selection pauses playback and seeks to the object.
- Drag inside moves. Side bars change width; top/bottom bars change height. The opposite edge stays anchored.
- Two fingers move, resize and freely rotate in one gesture. A dedicated corner control provides one-finger scale/rotate.
- Video has the same direct model on the Video tab, plus Fit, Fill, reset, numeric scale and a wide free-angle scrubber.
- Minimum touch target is 44–48 dp even when the visible handle is smaller. Continuous updates stay local; persistence happens at gesture end.
- Every destructive action is explicit and undoable. Delete remains available both on-object and in the inspector/timeline.

Orientation and canvas

Default canvas follows decoded source dimensions and rotation metadata. Switching to 9:16 never silently crops: Fit shows the whole clip; Fill covers the canvas and exposes pan/scale. Free rotation is independent of 90° source-orientation correction. Export must use the exact same normalized transform values as preview.

Technology and data architecture

Recommended Android stack

Layer	Choice	Reason
App shell	Expo SDK 57 / React Native 0.86 / Expo Router	Fast product iteration with native module escape hatches.
Preview	expo-video now; Media3/CompositionPlayer evaluation for export parity	Hardware-backed playback and an Android-native route to composition.
Transcription	whisper.rn 0.7 + whisper.cpp-compatible GGML models	No API key; model downloads enable offline use; MIT-licensed integration layer.
Audio decode	Small Kotlin Expo native module	Decode selected video to mono PCM WAV on device.
Projects	Expo SQLite + app files	Nondestructive persistence and resumable projects.
Export	Media3 Transformer + a tested caption compositing path	Android-supported editing/export primitives; parity and animated text need a prototype gate.
Fonts	Bundled families + user-imported TTF/OTF copied into project/app storage	Android system-font enumeration is inconsistent; import is controllable and testable.

Technical evidence: [whisper.cpp](#) • [whisper.rn](#) • [Media3 Transformer](#)

Style inheritance

RESOLUTION RULE

Resolved style = project default → caption override → word override. “All subtitles” changes the project default; “This subtitle” writes only a caption patch.

Layer	Stores	Example
Project	Full default style and animation configuration	Anton, white, 48, black stroke, active word lime.
Caption	Sparse override patch	Caption 7: yellow text, rotated 3°.
Word	Sparse override patch	FREE: green, 150%, punch.

Project invariants

- Source media is referenced and never rewritten during editing.
- Word timestamps remain canonical; caption blocks reference word IDs and can be regrouped.
- Coordinates and boxes are normalized to the output canvas so preview and export scale consistently.
- Video transform is independent of canvas preset: fit/fit, normalized position, scale and free rotation.

Commercial model

Recommended offer

Tier	Included	Price hypothesis
Core	Unlimited local captions, imports, fonts/presets, project editing, SRT/ASS, watermark-free 720p/1080p, no account.	Free and genuinely usable.
Supporter	Extra curated style/animation packs, batch export, advanced asset management, priority beta builds.	\$19–29 one-time during beta.
Studio later	Optional encrypted sync/backup, team libraries, collaboration, cloud fallback chosen by user.	\$4–8/month only when recurring service cost exists.

What must never become a paywall

- Transcription minutes or a required API key for ordinary local captions
- User-imported fonts, basic colors, timing correction, or This/All scope
- Saved presets, project count, export count, a watermark, or artificial resolution lock
- Access to a user's own project data or captions

Cost structure

Cost	Behavior	Control
Model delivery	Large one-time downloads can dominate bandwidth.	Offer tiny/base/small quantized choices, checksums, resumable downloads and optional mirrors.
Development/QA	Device fragmentation is the largest ongoing cost.	Test matrix by RAM, chipset, Android version, codec, resolution and thermal state.
Store economics	Fees vary by program, region and purchase model.	Treat Google's current service-fee policy as the source of truth; model scenarios, not one fixed percentage.
Support	Codec/export failures create expensive cases.	Collect an opt-in local diagnostic bundle with no media content.
Cloud later	Sync and cloud AI create real recurring cost.	Make them optional paid services and show when media leaves-device.

Commercial-policy references: [Google Play service fees](#) • [Target API requirements](#)

Trust, legal and platform

Privacy posture

Local-first can be a defensible promise only if the shipped app and every bundled SDK honor it. Video, extracted audio, transcripts and project data should stay on device unless the user explicitly chooses an external destination. Model downloads are network activity, but do not require uploading user media.

- Publish a short privacy policy even if no user data is collected off-device.
- Audit each SDK before release; analytics or crash tooling changes the Data safety declaration.
- Request photo/video permissions only at import/export time and explain the purpose in plain language.
- Provide Delete project and Delete downloaded models with storage estimates; never delete source media.
- If cloud features appear later, use an unmistakable Local / Cloud label and separate consent.

Google states that purely on-device access/processing is not treated as collected for the Data safety form, but every published app still needs an accurate form and privacy policy. [Google Play Data safety guidance](#)

Licensing checklist

Asset/code	Requirement	Gate
whisper.rn / whisper.cpp	Retain MIT notices and validate all transitive/native licenses.	Automated third-party notice inventory before store submission.
Whisper models	Document source, checksum, size, language behavior and redistribution terms.	Do not host a model until its license and bandwidth plan are recorded.
Fonts	Bundle only fonts with redistribution rights; imported fonts remain user-provided.	Per-font license file in the repository and in-app credits.
Icons/presets	Use original or properly licensed assets; avoid competitor trade dress.	Maintain source and provenance; no pixel-for-pixel Mojo imitation.
Export codecs	Rely on platform codecs where possible; review any bundled codec licenses.	Legal review before adding FFmpeg binaries.

Store readiness

The current build targets API 36 and supports Android 7+ at the manifest level. The sideloaded beta APK is arm64-v8a for modern phones. A Play release should use an Android App Bundle with ABI splits, signed release keys, crash-free testing, privacy policy, Data safety declaration, content rating, store assets and closed-track verification.

Go-to-market and measurement

Launch sequence

- 5. Private device alpha: Seeker and 3–5 contrasting Android devices; validate import, transcription, gestures, orientation, backgrounding and thermal behavior.
- 6. Creator design partners: 10–20 users who publish weekly; require them to finish real clips, not explore a demo.
- 7. GitHub beta: signed APK, checksums, clear model sizes, issue template and privacy promise.
- 8. Play closed test: universal AAB, policy declarations, codec matrix and export reliability gate.
- 9. Public launch: short before/after clips demonstrating local mode, one/all scope, unified fonts and TikTok fit/fill—not a generic feature montage.

North-star and guardrails

Metric	Definition	Initial gate
Captioned export success	% of imported projects that produce a playable exported MP4 within 24h.	≥55% beta; ≥70% launch.
Time to first editable caption	Import tap to generated caption selection, excluding intentional model download.	Median ≤ video duration × 1.5 on Balanced target hardware.
Correction burden	Edited words ÷ transcribed words for completed projects.	Track by language/device; reduce without hiding uncertainty.
Gesture success	% of move/resize/rotate attempts completed without immediate undo/retry.	≥90% in observed task tests.
Export reliability	% of export jobs producing valid A/V sync and intended orientation.	≥98% across release codec matrix.
Privacy guardrail	Unexpected media/transcript bytes sent off-device.	Zero in local mode.

Distribution message

POSITIONING

The private, no-limits caption editor for Android: fast local transcription, TikTok-ready framing, and precise typography without subscriptions getting between you and your own video.

Risk register and roadmap

Risk	Impact	Mitigation / release gate
Preview/export mismatch	Critical	One normalized scene model; golden-frame tests at timestamps; ship no animation until parity passes.
Touch gestures still feel poor	Critical	Observed one-hand tests, 48 dp hit zones, gesture-end persistence, undo, instrumentation without media content.
Transcription too slow/hot	High	Quantized model choices, benchmark matrix, thermal cancellation, display storage/time estimates.
Codec/device fragmentation	High	Media3 compatibility matrix, proxy preview, fallback error messages, sample diagnostic files.
Model bandwidth	Medium	Tiny default, opt-in larger models, resumable verified downloads, delete controls.
Scope creep into CapCut	High	No broad effects until captioned export and correction metrics meet gates.
Open-source/commercial licensing	High	Third-party notices, font provenance, no bundled codec without review.
“Local-first” trust breach	Critical	No trackers by default; network audit; explicit cloud mode later.

Recommended delivery order

Phase	Deliverable	Exit criterion
0 — Acceptance	Seeker gesture/orientation test of current beta.	Move, edge resize, pinch/rotate, video scale and timing/delete pass observed tasks.
1 — Pipeline	Import → local words → grouped captions → edit → burned MP4 + SRT.	Ten real files export with correct A/V sync and orientation.
2 — Serious editor	Fonts/import, one/all, split/merge, timing, background/stroke/shadow/position.	Creators finish clips without leaving the app.
3 — Motion	Active word, karaoke, pop, bounce, punch, unlimited data-driven presets.	Preview/export golden-frame parity.
4 — Power	Per-word overrides, advanced type and animation packs.	Schema remains backward compatible.
5 — Video	Trim, split, crop, speed, audio, fades, safe zones.	Adds retention without slowing core caption path.

Appendix

Current artifact

Field	Value
App	Caption Studio 1.0.0 beta
Package	app.captionstudio.mobile
Release APK	caption-studio-android.apk
APK ABI	arm64-v8a
Minimum / target SDK	24 / 36
Architecture	Expo SDK 57, React Native 0.86, Kotlin native audio module, whisper.rn 0.7, SQLite
Source behavior	Nondestructive project edits; original media is never modified by editor operations

Source register

- [Mojo Pro](https://mojo-app.com/pro/) — <https://mojo-app.com/pro/>
- [Mojo Pro AI help](https://help.mojo-app.com/en/articles/11036576-subscribing-to-mojo-pro-ai) — <https://help.mojo-app.com/en/articles/11036576-subscribing-to-mojo-pro-ai>
- [CapCut pricing](https://www.capcut.com/help/how-much-does-capcut-pro-cost) — <https://www.capcut.com/help/how-much-does-capcut-pro-cost>
- [CapCut Standard vs Pro](https://www.capcut.com/help/capcut-standard-vs-pro) — <https://www.capcut.com/help/capcut-standard-vs-pro>
- [Captions plans](https://www.captions.ai/plans) — <https://www.captions.ai/plans>
- [Captions Android](https://www.captions.ai/download) — <https://www.captions.ai/download>
- [Google Play target API](https://developer.android.com/google/play/requirements/target-sdk) — <https://developer.android.com/google/play/requirements/target-sdk>
- [Google Play service fees](https://support.google.com/googleplay/android-developer/answer/112622?hl=en-CA) — <https://support.google.com/googleplay/android-developer/answer/112622?hl=en-CA>
- [Google Play Data safety](https://support.google.com/googleplay/android-developer/answer/10787469?hl=en) — <https://support.google.com/googleplay/android-developer/answer/10787469?hl=en>
- [Android Media3 Transformer](https://developer.android.com/media/media3/transformer/getting-started) — <https://developer.android.com/media/media3/transformer/getting-started>
- [whisper.cpp](https://github.com/ggml-org/whisper.cpp) — <https://github.com/ggml-org/whisper.cpp>
- [whisper.rn](https://github.com/mybigday/whisper.rn) — <https://github.com/mybigday/whisper.rn>

Research cutoff: 16 August 2026. Published pricing and platform rules can change; re-check primary sources immediately before a commercial or store decision.